

Implementing the ls command

Module 5 self study material

Operating systems 2018

1DT044 and 1DT096



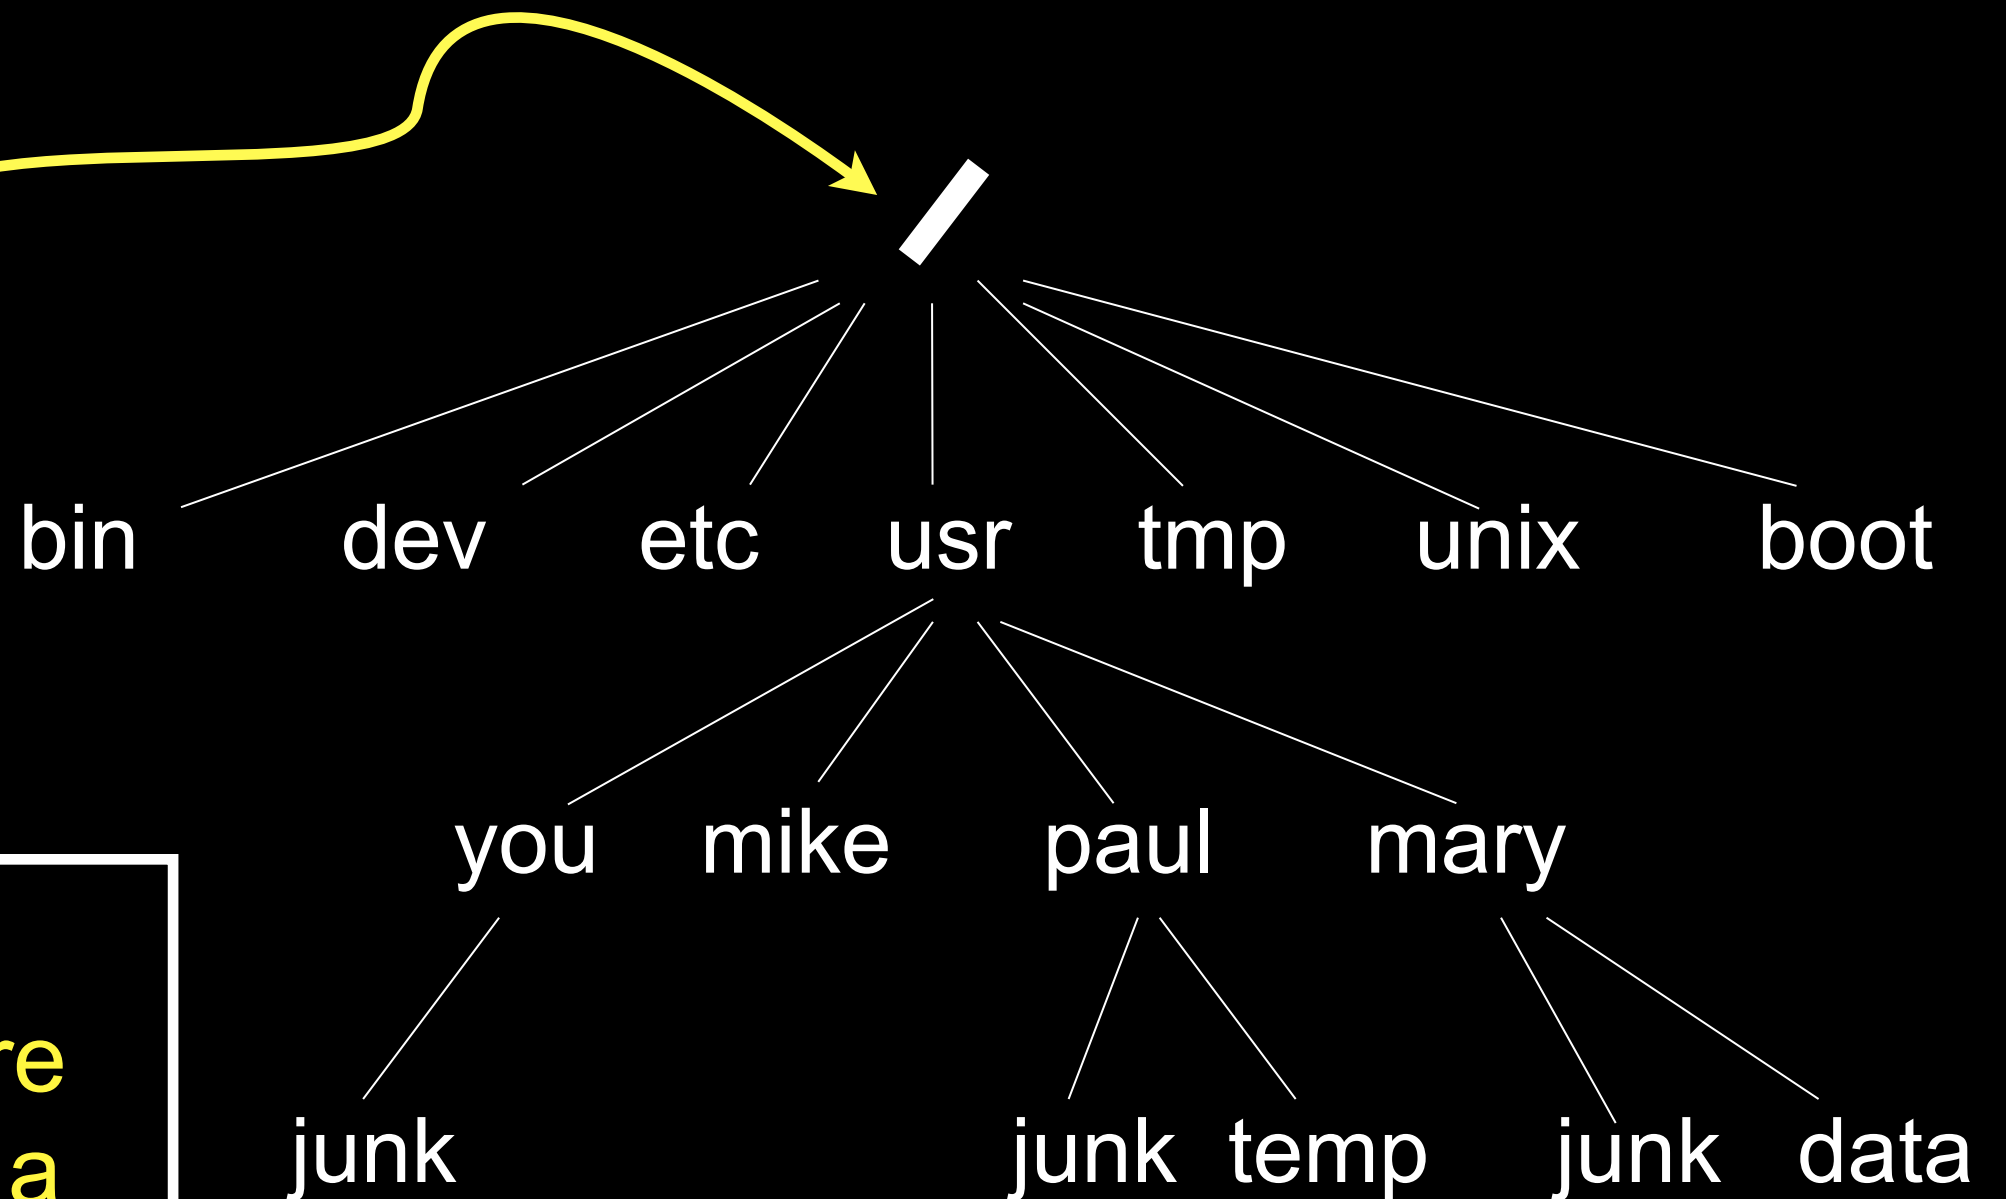
The Unix command interpreter (aka shell) does not understand the commands in any way.

Commands are implemented through system programs.

The shell creates a new process using `fork()` and have the new process use `exec()` to run the identified system program executable.

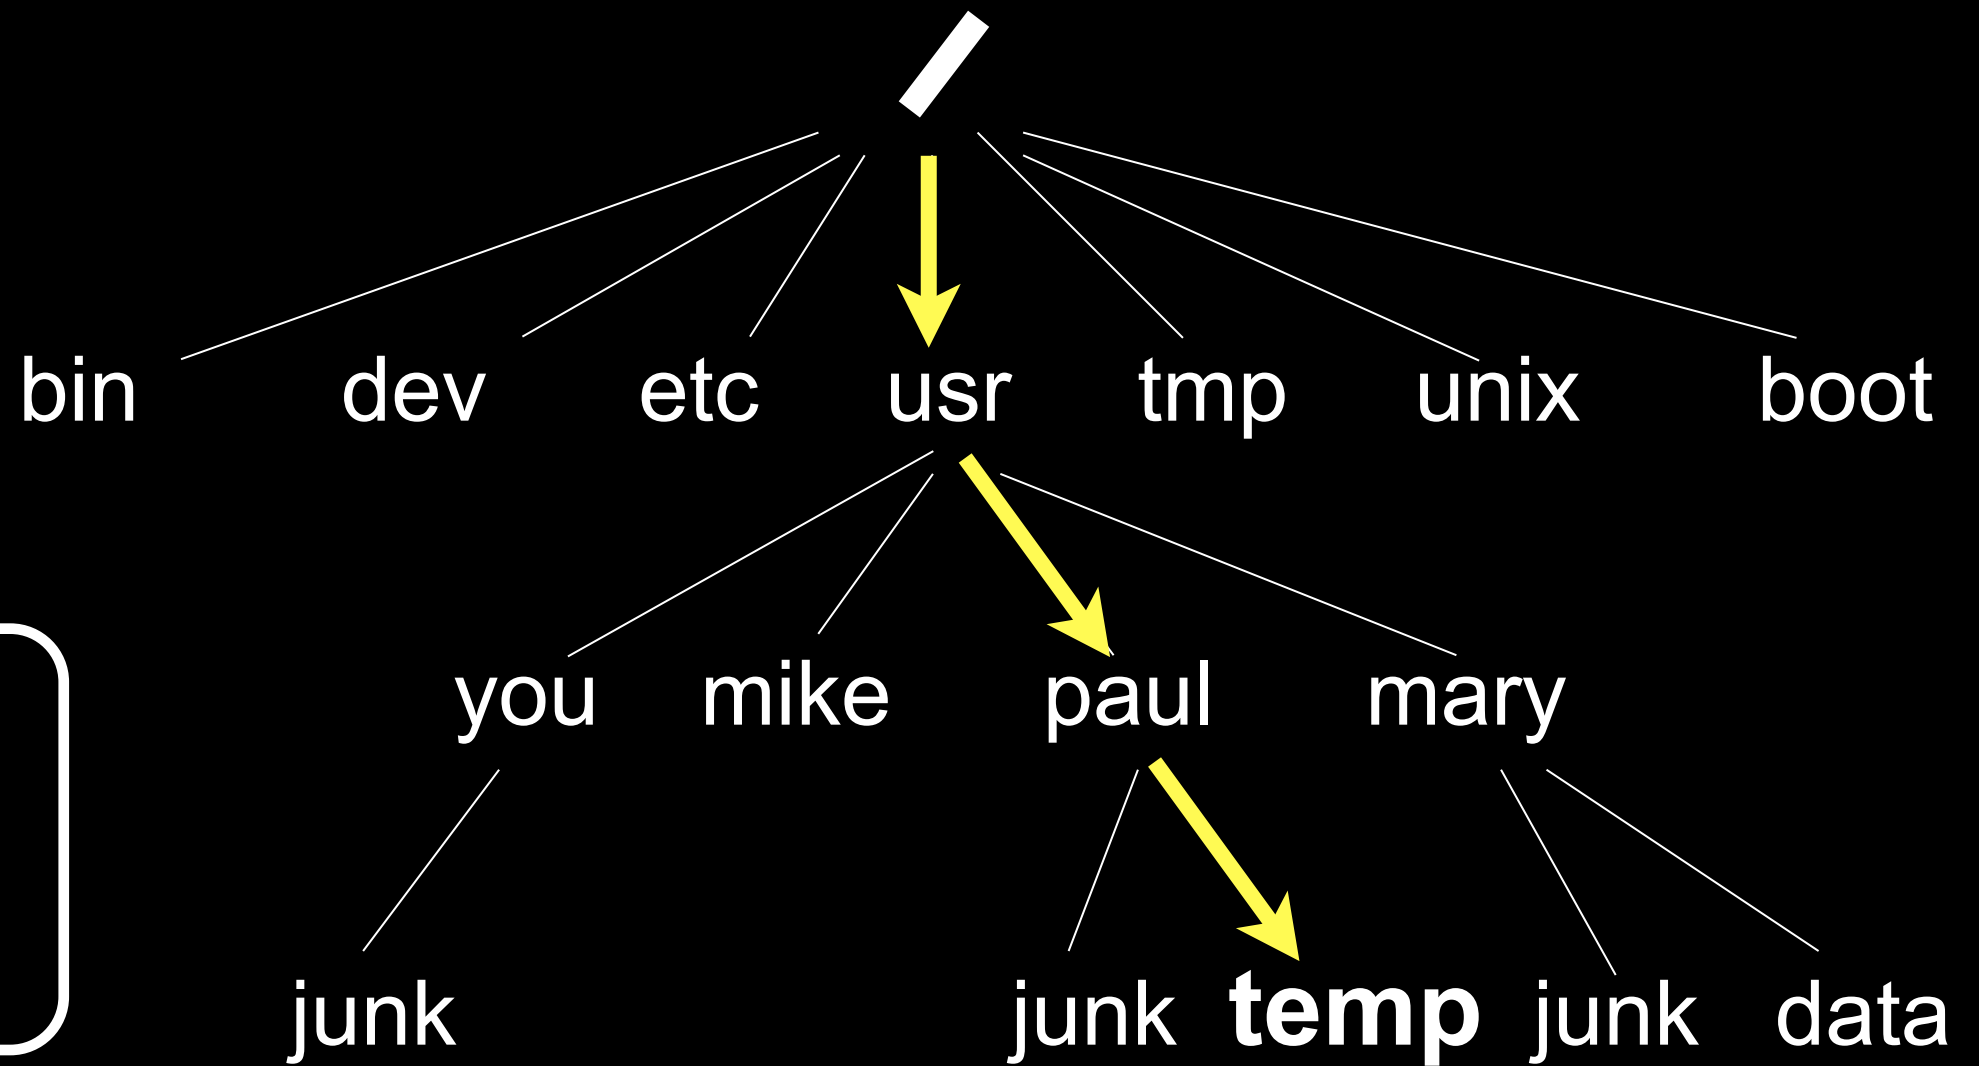
We should be able to implement our own (simplified) version of the `ls` command.

The Unix file system



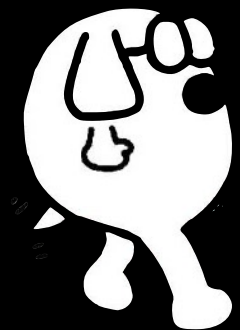
Files and directories are organized in a **tree structure** – starting at the **root**.

The Unix file system

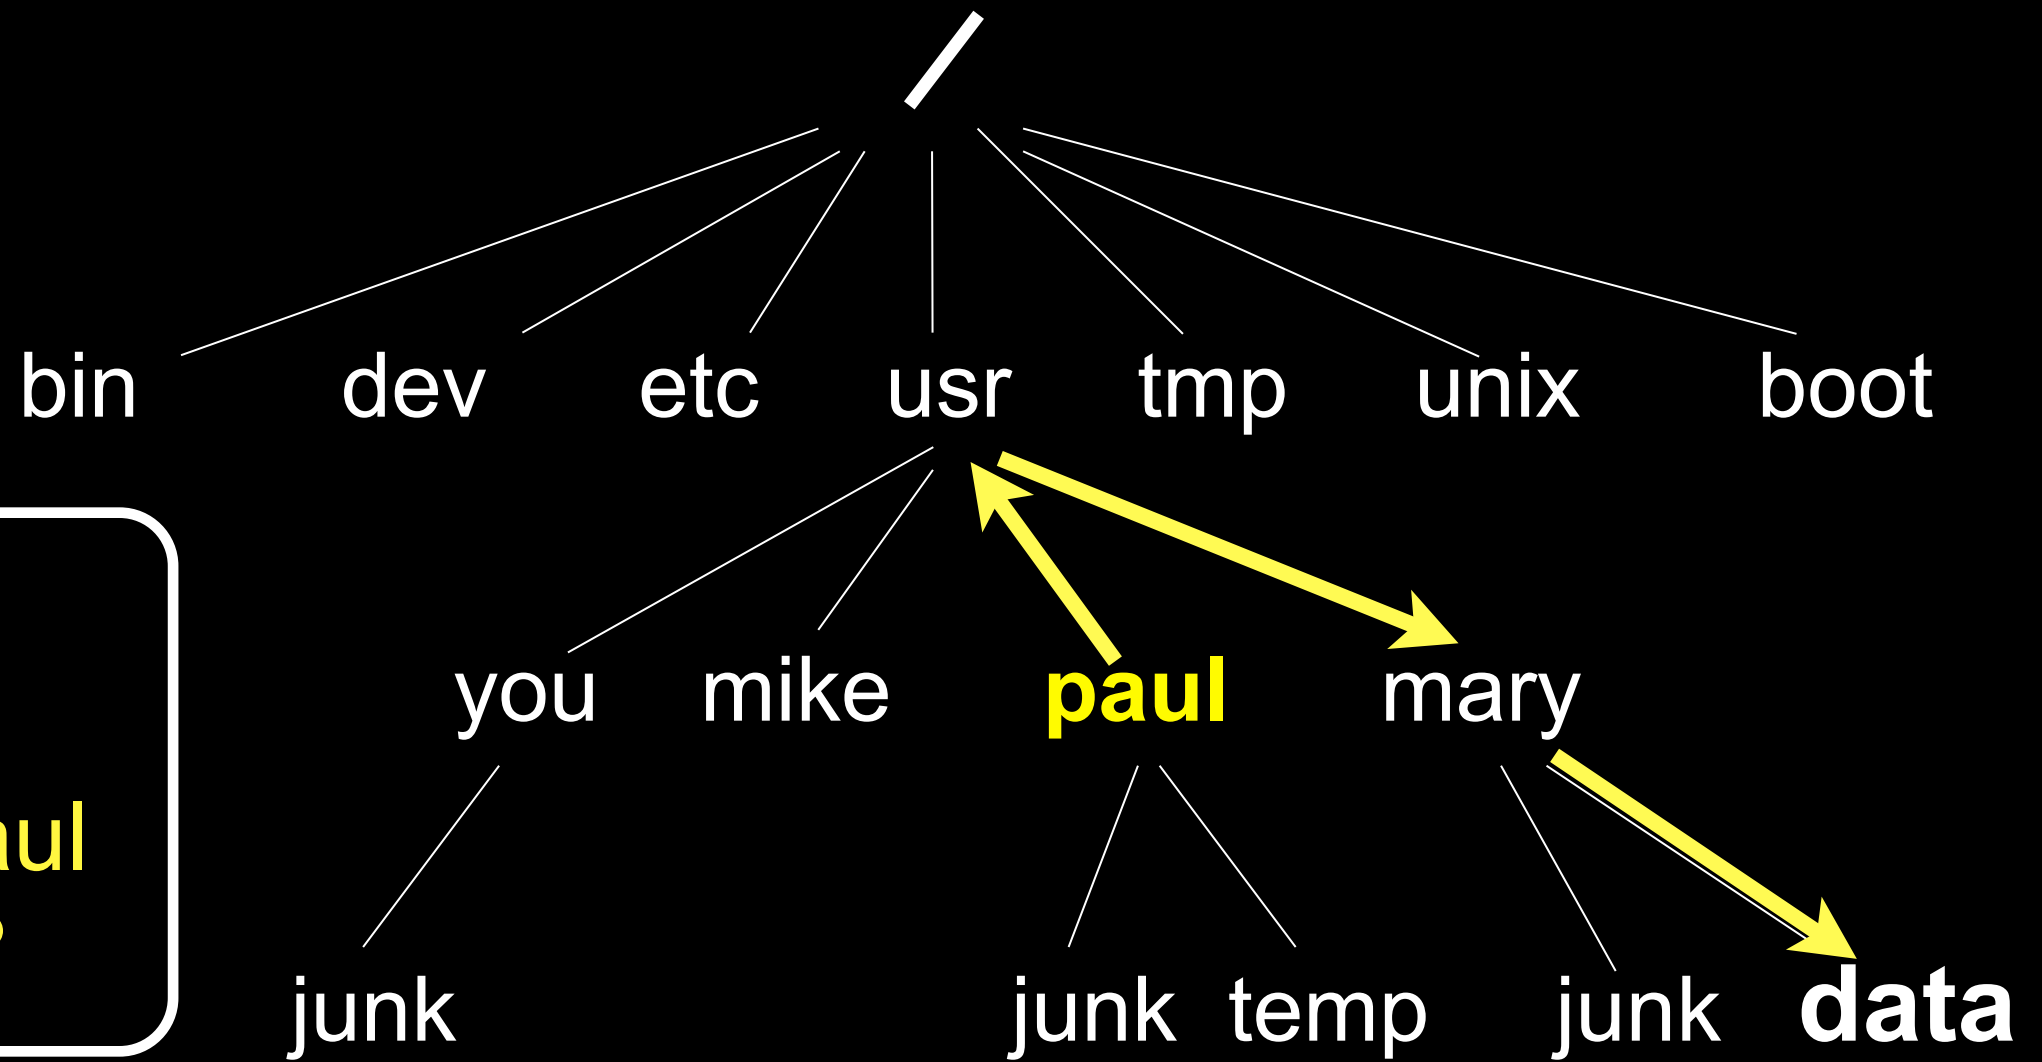


**Absolute
path to file
temp?**

/usr/paul/temp

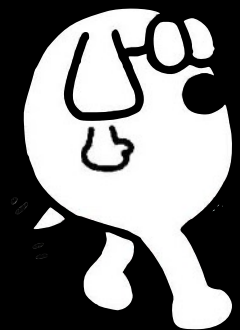


The Unix file system



Relative
path from
directory paul
to file data?

../mary/data



The **ls** command

In Unix-like systems you can use the ls command to get information about files in a directory.

The **owner** of the file

Highest **group** that the owner belongs to

```
$> ls -l
```

```
-rw----- 1 hans it readme.txt
```

```
-rwx-r--r-- 1 karl it script.sh
```

File mode - **permissions**

Name of File or Directory

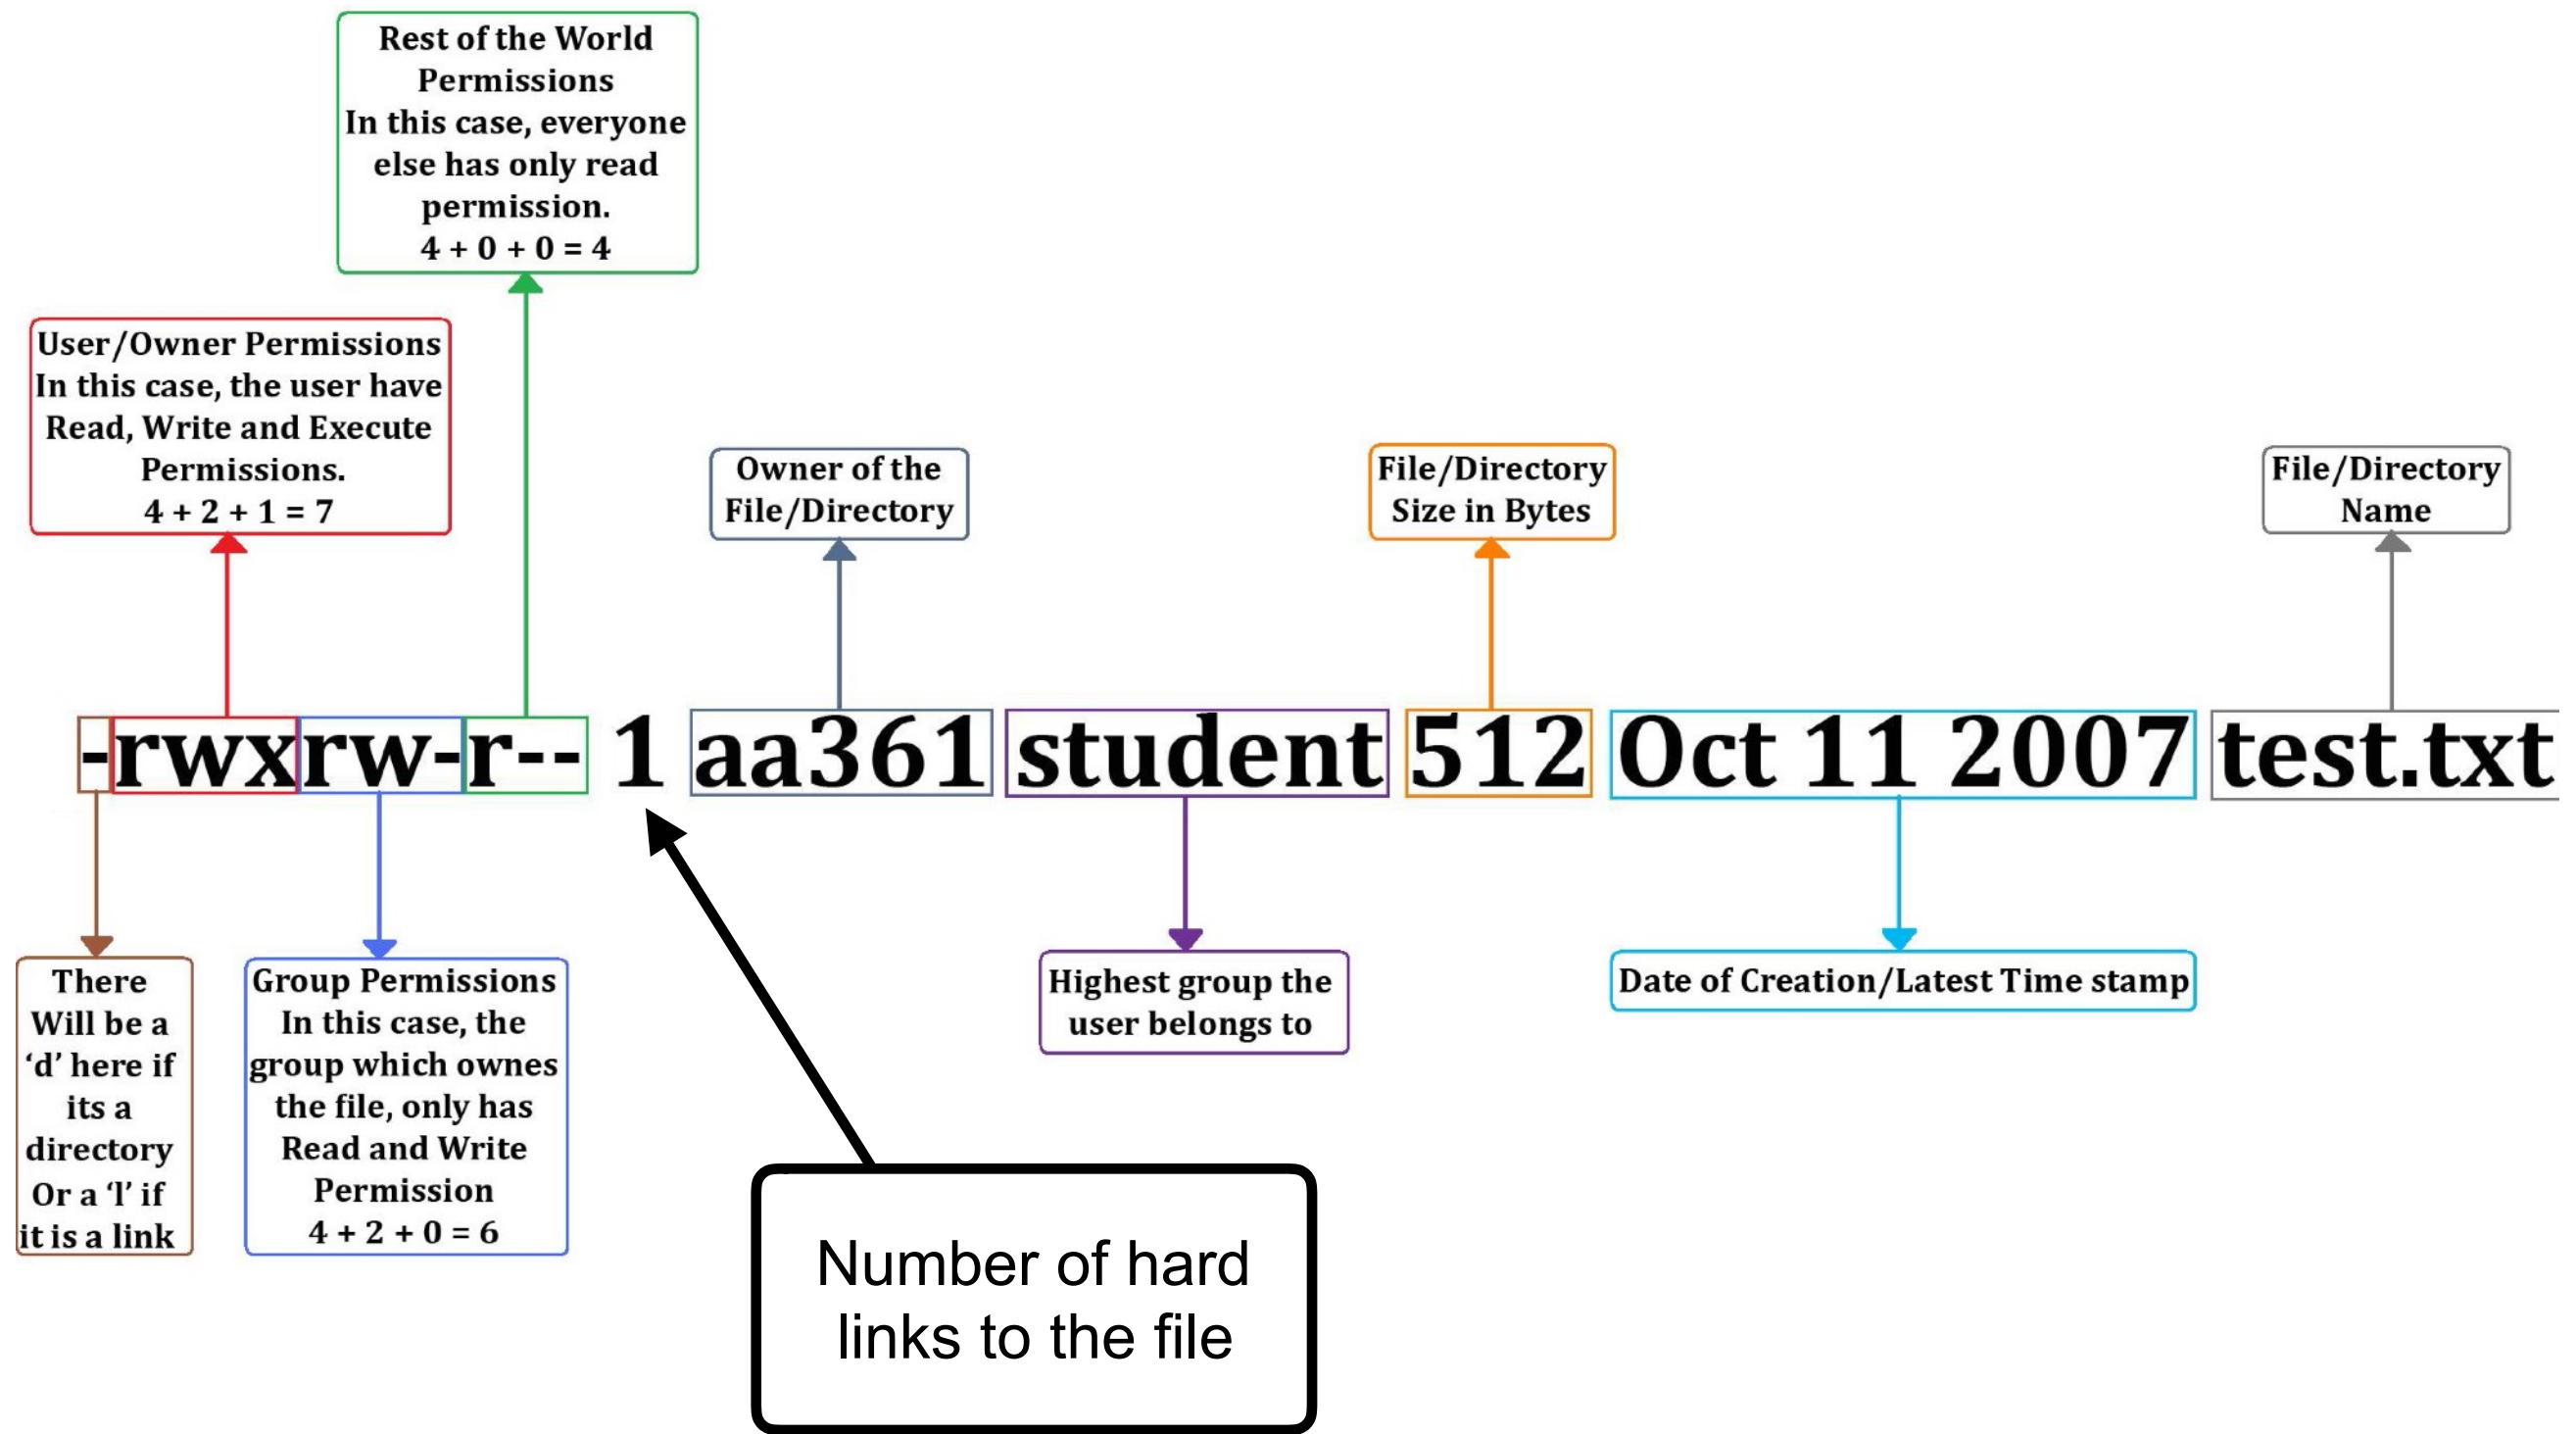
Unix file permissions

```
$> ls -l
```

```
-rw----- 1 hans it readme.txt  
-rwx-r--r-- 1 karl it script.sh
```

r = read **w** = write **x** = execute

- r w x r - - r - -
 └──┬──┘ └──┬──┘ └──┬──┘
 owner group all
 others



In Unix-like system a file is represented by exactly one **inode**.



inode

Examples of information stored in an inode

File mode

Size

inode number

Reference count

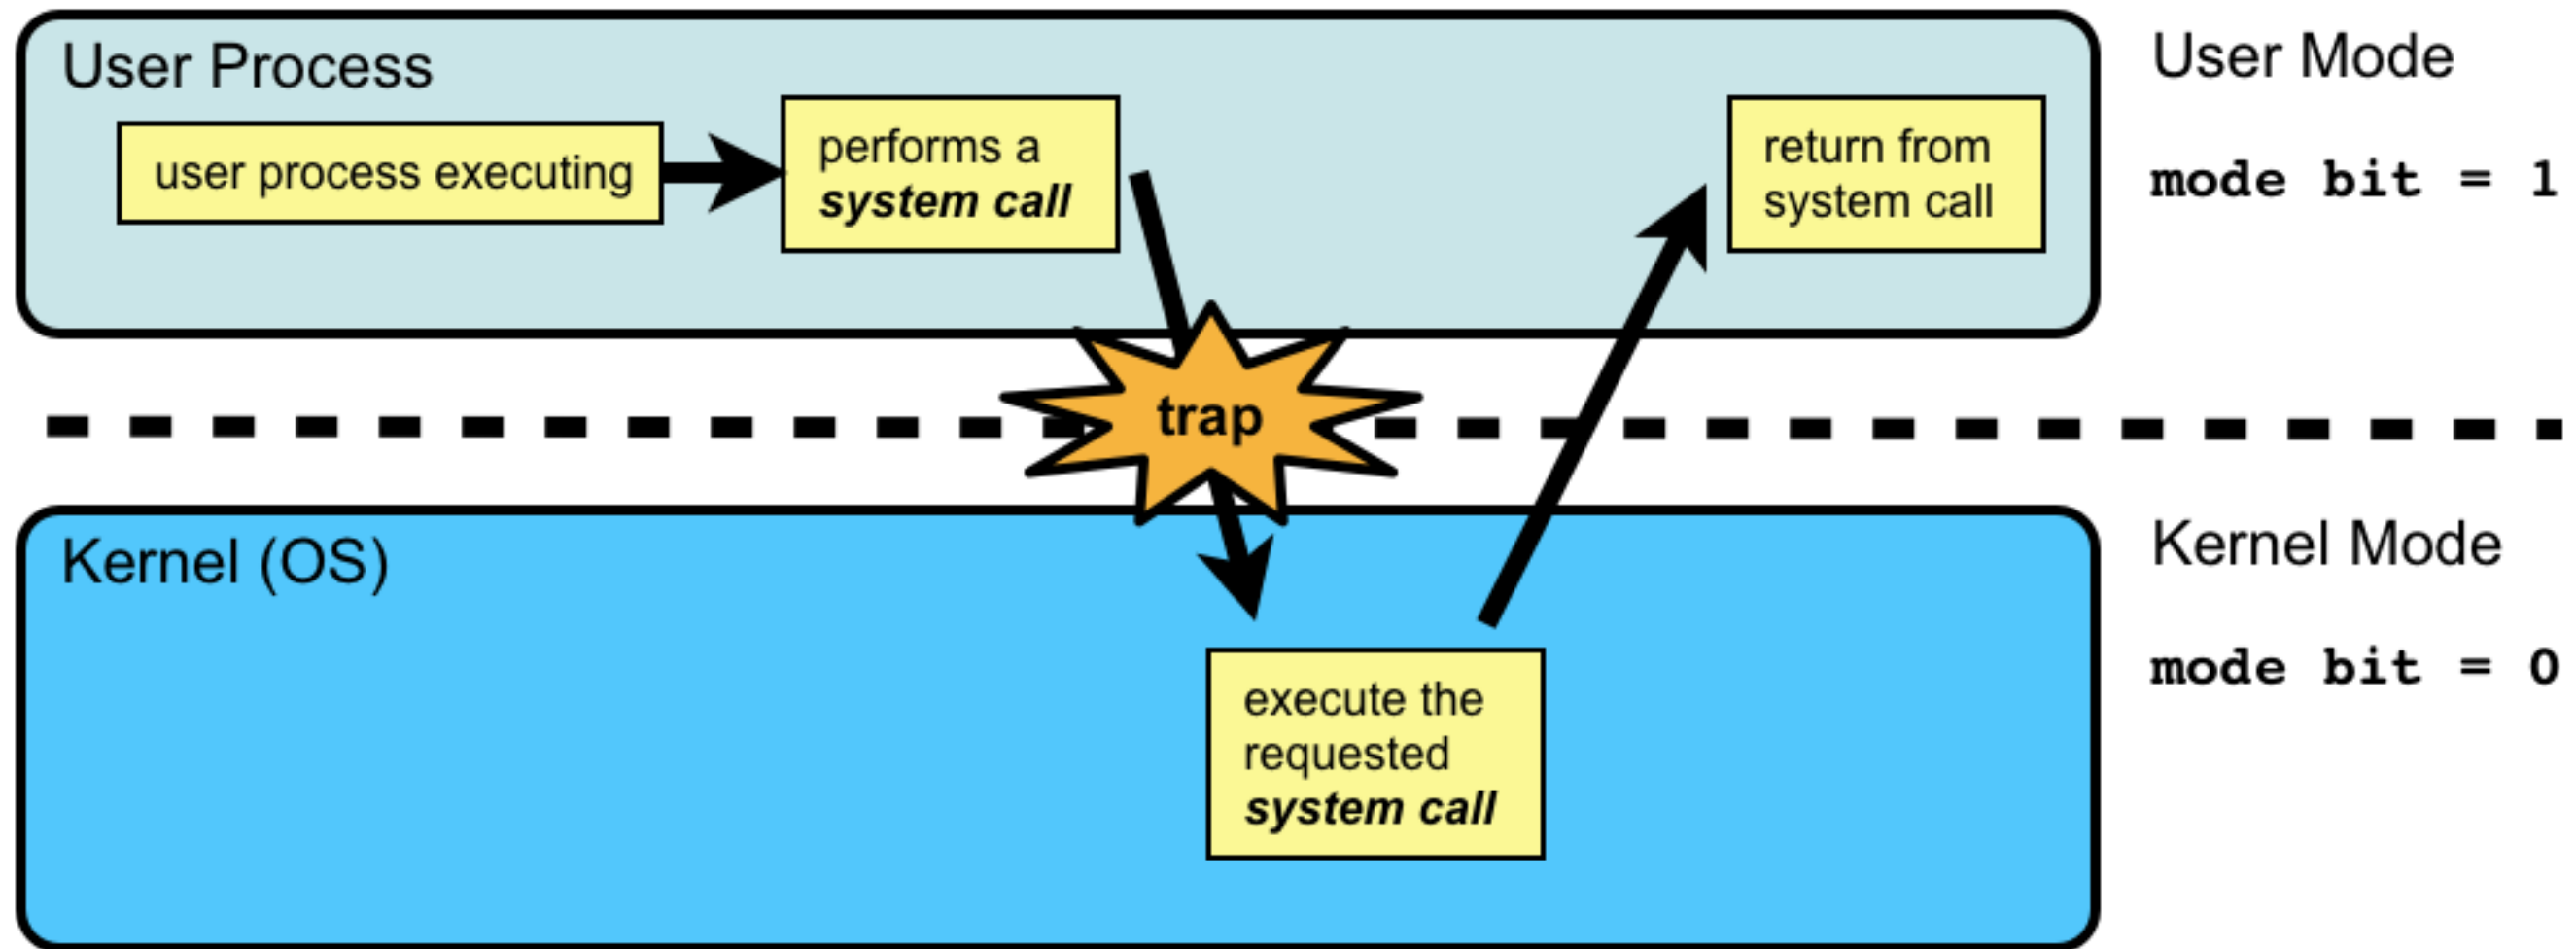
Owner

Pointer to data

Timestamps

- ▶ A file does not have a name.
- ▶ The file is uniquely identified by its inode number.
- ▶ The "file name" is a property of the directory.

To obtain information about files and directories there exist a number **system calls**.



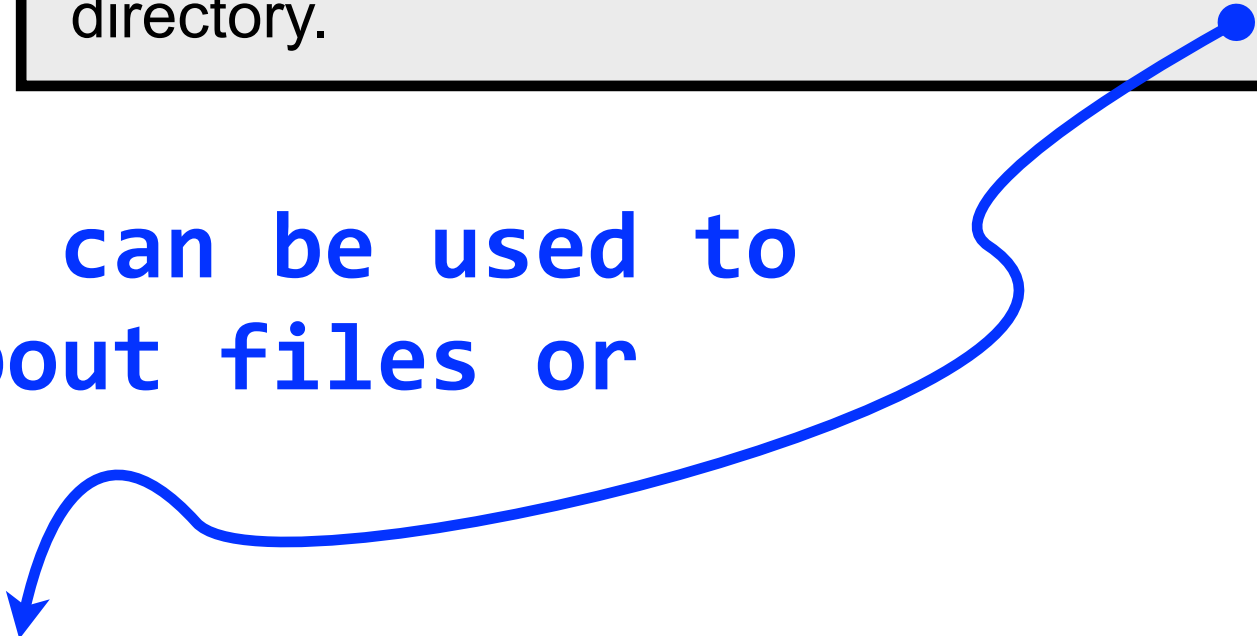
C programming

with

stat() & lstat()

```
#include <sys/types.h>
#include <sys/stat.h>
```


Relative or absolute path to file or directory.



```
// Two system calls that can be used to
// obtain information about files or
// directories.
```

```
int  stat(const char *path, struct stat *buf);
int  lstat(const char *path, struct stat *buf);
```

```
// Both stat() & lstat() returns 0 on success
// and -1 on failure.
```



Must pass a **pointer** to a **stat struct** – will hold result of the system call.

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
```

```
struct stat fstat;
int status;
```

```
status = stat("./file.txt", &fstat);
```

```
printf("MODE = %d \n", (int) fstat.st_mode);
```

```
// Convert the numeric mode to the standard
// string representation such as "-rw-r--r--".
```

```
printf("MODE = %s \n", strmode(fstat.st_mode));
```

Don't forget to use the address-of operator & to get the pointer to **fstat**.

Now the **fstat** struct will be populated with data.

Need to implement a **strmode()** function that converts the numeric mode to string.

See the man page for the stat struct for more information

```
$ man -s2 stat
```

```
struct stat {
    dev_t      st_dev;          /* ID of device containing file */
    ino_t      st_ino;          /* inode number */
    mode_t     st_mode;         /* protection */
    nlink_t    st_nlink;        /* number of hard links */
    uid_t      st_uid;          /* user ID of owner */
    gid_t      st_gid;          /* group ID of owner */
    dev_t      st_rdev;         /* device ID (if special file) */
    off_t      st_size;         /* total size, in bytes */
    blksize_t  st_blksize;      /* blocksize for filesystem I/O */
    blkcnt_t   st_blocks;       /* number of 512B blocks allocated */
    .
    .
    .
};
```

User name

How can we get the user name (string) of the file owner?

```
$> ls -l
```


```
-rw----- 1 hans it readme.txt
```

```
-rwx-r--r-- 1 karl it script.sh
```

See the man page for the stat struct for more information

```
$ man -s2 stat
```

```
struct stat {
    dev_t      st_dev;          /* ID of device containing file */
    ino_t      st_ino;          /* inode number */
    mode_t     st_mode;         /* protection */
    nlink_t    st_nlink;        /* number of hard links */
    uid_t      st_uid;          /* user ID of owner */
    gid_t      st_gid;          /* group ID of owner */
    dev_t      st_rdev;         /* device ID (if special file) */
    off_t      st_size;         /* total size, in bytes */
    blksize_t  st_blksize;      /* blocksize for filesystem I/O */
    blkcnt_t   st_blocks;       /* number of 512B blocks allocated */
    .
    .
    .
};
```



getpwnuid()

Convert a numerical user id to username string.

The `getpwnuid()` function shall return a pointer to a `struct passwd` with the structure as defined in `<pwd.h>` with a matching entry if found.

```
struct passwd *getpwnuid(uid_t uid);
```

pwd.h

The `<pwd.h>` header shall provide a definition for **struct passwd**, which shall include at least the following members:

<code>char</code>	<code>*pw_name</code>	<code>// User's login name.</code>
<code>uid_t</code>	<code>pw_uid</code>	<code>// Numerical user ID.</code>
<code>gid_t</code>	<code>pw_gid</code>	<code>// Numerical group ID.</code>
<code>char</code>	<code>*pw_dir</code>	<code>// Initial working directory.</code>
<code>char</code>	<code>*pw_shell</code>	<code>// Program to use as shell.</code>

Get user name of the file owner

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <pwd.h>
```

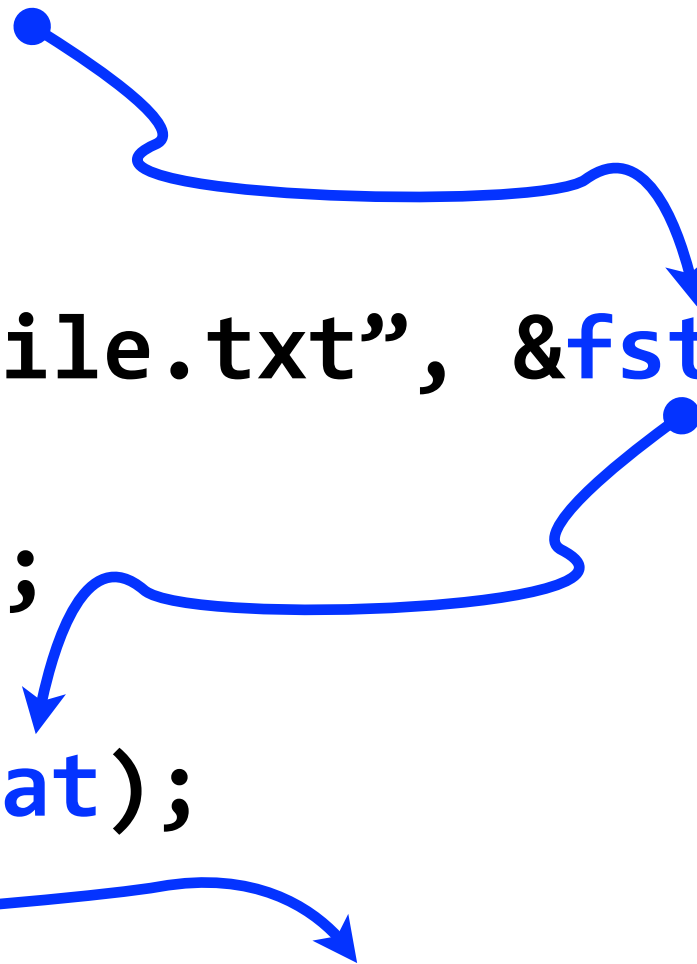
```
struct stat fstat;
int status;
```

```
status = stat("./file.txt", &fstat);
```

```
struct passwd *pwd;
```

```
pwd = getpwuid(fstat);
```

```
printf("USER = %s \n", pwd->pw_name);
```



Obtaining information about links

```
$> ls -l
```

```
lrw----- 1 karl it link -> file.txt  
-rw-r--r-- 1 karl it file.txt
```

- ▶ Directory entry **link** is a symbolic link to **file.txt**.
- ▶ How to find out if something is a symbolic link?
- ▶ For a symbolic link, to what does the link refer?

The difference between `stat()` & `lstat()`

```
int  stat(const char *path, struct stat *buf);
```

```
int  lstat(const char *path, struct stat *buf);
```

- ▶ The `lstat()` function shall be equivalent to `stat()`, except when path refers to a symbolic link.
- ▶ In that case `lstat()` shall return information about the link, while `stat()` shall return information about the file the link references.

Obtaining information about links

```
#define MAX_LINK_LEN 20
```

```
struct stat fstat;  
lstat("link", &fstat);
```

Use **lstat()** so we can get info about symbolic links.

S_ISLNK(m) macro that tests whether file with mode **m** is a symbolic link or not.

```
// Using lstat() we can check if a file  
// is a symbolic link or not:
```

```
if (S_ISLNK(fstat.st_mode)) {  
    char link_buffer[MAX_LINK_LEN];  
    int len;  
    len = readlink("link", link_buffer, MAX_LINK_LEN);  
    link_buffer[len] = 0; // Null terminate  
    printf("link -> %s\n", link_buffer);  
}
```

readlink() system call used to get contents of symbolic link.